

Avalanche: metastable gossip consensus

Péter Garamvölgyi


outline

- **what is metastability?**
- the Snow consensus family
 - Slush
 - Snowflake
 - Snowball
- Avalanche

what is metastability?

*“Metastable state is a concept in physics [...] particles may be in a **mixed state**, but if properly induced, they can **quickly stabilize**, that is, all particles share the same state [...] In this state, the particles phase shift to form an invariable and stable structure. It's a dynamic chaotic process that stabilizes steadily, which is exactly what we want to happen in distributed systems.” **

what is metastability?



1: metastable state

2: unstable state

3: stable state

metastability in binary consensus

- nodes need to decide between **RED** and **BLUE**
 - **50%–50%**: unstable (bivalent) state
 - **51%–49%**: metastable state
 - **99%– 1%**: stable state
- goal: quick and robust convergence to an irreversible stable state

outline

- what is metastability?
- the Snow consensus family
 - **Slush**
 - Snowflake
 - Snowball
- Avalanche

Slush: non-BFT metastable binary consensus

```
def onQuery(v, col):  
    if color == None: color = col  
    respond(v, color)  
  
def slush(me, col0):  
    color = col0  
  
    for _ in range(m):  
        if color == None: return None  
  
        K = sample(N \ {me}, k)  
        P = map(lambda v: query(v, color), K)  
  
        for c in {RED, BLUE}:  
            if P.count(c) >= alpha: color = c  
  
    return color
```

main idea:

iterative random sampling

parameters:

N: full (?) view of network participants
m: number of rounds
k: sample size
alpha: majority parameter ($\alpha > k/2$)

Slush: non-BFT metastable binary consensus

```
def onQuery(v, col):
    if color == None: color = col
    respond(v, color)

def slush(me, col0):
    color = col0

    for _ in range(m):
        if color == None: return None

        K = sample(N \ {me}, k)
        P = map(lambda v: query(v, color), K)

        for c in {RED, BLUE}:
            if P.count(c) >= alpha: color = c

    return color
```

properties

almost memoryless; no history

$O(nk)$ communication overhead ($k \ll n$)

random sampling breaks 50/50 ties

Slush: irreversibility

```
def onQuery(v, col):
    if color == None: color = col
    respond(v, color)

def slush(me, col0):
    color = col0

    for _ in range(m):
        if color == None: return None

        K = sample(N \ {me}, k)
        P = map(lambda v: query(v, color), K)

        for c in {RED, BLUE}:
            if P.count(c) >= alpha: color = c

    return color
```

model using Continuous-Time Markow Chains

state at time t : $\mathcal{S}_t = n/2 + \delta$

i.e. δ more **BLUE** than **RED**

prob of reverting to **minority value** bounded by

$$\xi_\delta \leq \left(\frac{1/2 - \delta/n}{\alpha/k} \right)^\alpha \left(\frac{1/2 + \delta/n}{1 - \alpha/k} \right)^{k-\alpha} \\ \leq e^{-2((\alpha/k) - (1/2) + (\delta/n))^2 k}$$

... drops exponentially with δ

... can be arbitrarily small by tuning k and α

Slush: safety

```
def onQuery(v, col):  
    if color == None: color = col  
    respond(v, color)  
  
def slush(me, col0):  
    color = col0  
  
    for _ in range(m):  
        if color == None: return None  
  
        K = sample(N \ {me}, k)  
        P = map(lambda v: query(v, color), K)  
  
        for c in {RED, BLUE}:  
            if P.count(c) >= alpha: color = c  
  
    return color
```

what about Byzantine nodes?

when nodes develop preference for one color

... adversaries can flip nodes to the opposite

... keeping the network in balance

... and preventing consensus.

outline

- what is metastability?
- the Snow consensus family
 - Slush
 - **Snowflake**
 - Snowball
- Avalanche

Snowflake: adding BFT

```
def snowflake(me, col0):
    color = col0
    count = 0

    while True:
        if color == None: return None
        K = sample(N \ {me}, k)
        P = map(lambda v: query(v, color), K)

        for c in {RED, BLUE}:
            if P.count(c) >= alpha:
                count = 0 if color != c else count+1
                color = c

        if count > beta: return color
```

new idea:

nodes must explicitly detect irreversibility

capture strength of conviction using a counter

decide after *beta* identical consecutive samples

new parameters:

beta: decision threshold


Snowflake: safety argument

```
def snowflake(me, col0):  
    color = col0  
    count = 0  
  
    while True:  
        if color == None: return None  
        K = sample(N \ {me}, k)  
        P = map(lambda v: query(v, color), K)  
  
        for c in {RED, BLUE}:  
            if P.count(c) >= alpha:  
                count = 0 if color != c else count+1  
                color = c  
  
        if count > beta: return color
```

A: node a decides on **BLUE** at t_1

B: node b decides on **RED** at t_2

failure probability: $p(\mathbf{B} \mid \mathbf{A}) \leq \epsilon$



outline

- what is metastability?
- the Snow consensus family
 - Slush
 - Snowflake
 - **Snowball**
- Avalanche

Snowball: reduce random perturbations

```
def snowball(me, col0):
    color = lastc = col0
    conf[RED] = conf[BLUE] = 0
    count = 0

    while True:
        if color == None: return None
        K = sample(N \ {u}, k)
        P = map(lambda v: query(v, color), K)

        for c in {RED, BLUE}:
            if P.count(c) >= alpha:
                count = 0 if c != lastc else count+1
                lastc = c

                conf[c]++
                if conf[c] > conf[color]: color = c

        if count > beta: return color
```

new idea:




introduce history (confidence)

only change preference based on total confidence

new parameters:

beta: decision threshold

Snowball



outline

- what is metastability?
- the Snow consensus family
 - Slush
 - Snowflake
 - Snowball
- **Avalanche**

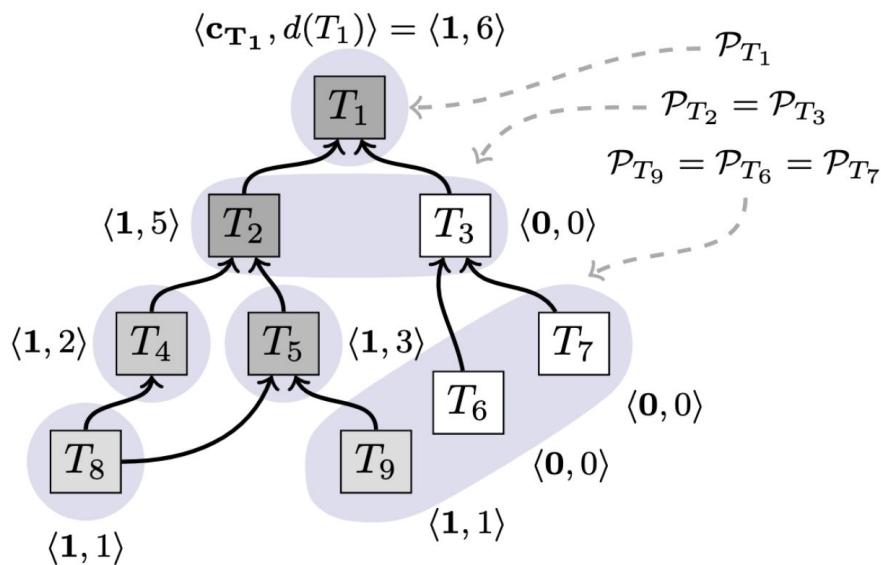
Avalanche: DAG-based digital payments

- observation: correct clients never double spend (attackers cannot forge signatures)
 - ⇒ safety and liveness guaranteed for virtuous transactions
 - ⇒ no liveness guarantee for rogue transactions

Avalanche: DAG-based digital payments


- maintain append-only DAG of transactions
 - application-defined parent-child relationship (not the same as UTXO!)
 - vertices belong to conflict sets
 - one Snowball instance for each conflict set → choose one tx
- why DAG?
 - a vote on a vertex implicitly votes for all ancestors
 - confidence is derived from votes on predecessors
 - past decisions are harder to undo

Avalanche: DAG-based digital payments



- vertices collect (immutable) 0/1 **chit** values using a one-time query
- **confidence** is derived from sub-DAG and grows as the DAG grows
- a vertex is **strongly preferred** if all its ancestors are the preferred one in their respective conflict sets

Avalanche: multi-input UTXOs



- financial transactions are embedded into Avalanche vertices, i.e. we have 2 DAGs
- each input corresponds to a single vertex
- tx accepted if all inputs are accepted

references

Scalable and Probabilistic Leaderless BFT Consensus through Metastability. (2019).

<https://avalabs.org/snow-avalanche.pdf>