

NS 2018 Paper Reading

Péter Garamvölgyi 2018280070

Outline

- **GPU-based packet classification [1]**
 - packet classification overview
 - problem statement
 - GPU programming
 - bloom filters
 - GSwitch
 - results & evaluation
- safe network updates [3]
 - problem statement
 - inconsistent update example
 - update mechanisms
 - mathematical model
 - results & evaluation

packet classification - overview

given a rule-/filter-set...

... and a packet (usually characterized by the standard 5-tuple) ...

... decide which rules match ...

... and execute the associated actions.

packet classification - example filter set

<i>Filter</i>				<i>Action</i>	
<i>SA</i>	<i>DA</i>	<i>Prot</i>	<i>DP</i>	<i>FlowID</i>	<i>PT</i>
11010010	*	TCP	[3:15]	0	3
10011100	*	*	[1:1]	1	5
101101*	001110*	*	[0:15]	2	8†
10011100	01101010	UDP	[5:5]	3	2
*	*	ICMP	[0:15]	4	9†
100111*	011010*	*	[3:15]	5	6†
10010011	*	TCP	[3:15]	6	3
*	*	UDP	[3:15]	7	9†
11101100	01111010	*	[0:15]	8	2
111010*	01011000	UDP	[6:6]	9	2
100110*	11011000	UDP	[0:15]	10	2
010110*	11011000	UDP	[0:15]	11	2
01110010	*	TCP	[3:15]	12	4†
10011100	01101010	TCP	[0:1]	13	3
01110010	*	*	[3:3]	14	3
100111*	011010*	UDP	[1:1]	15	4

problem statement

- virtualization and SDN made software switches popular again
- software-based packet classification has unique challenges
 - large, dynamic rule tables
 - multidimensional tuples
- packet classification has become the bottleneck

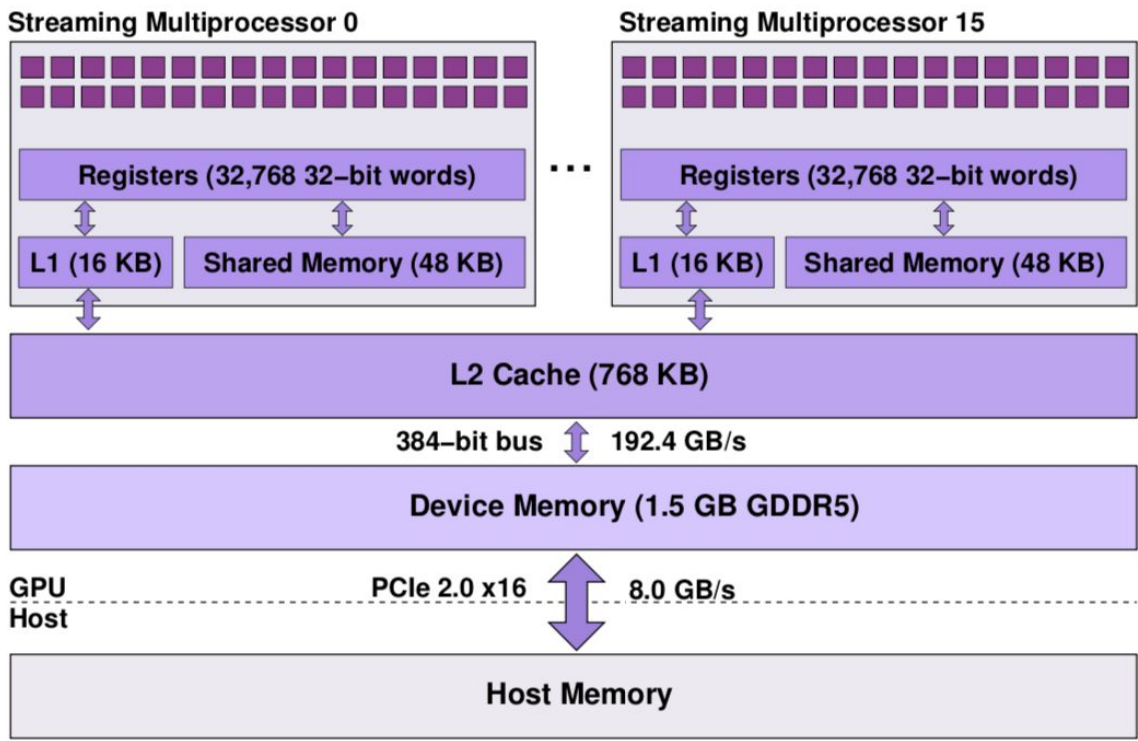
paper contributions

1. GPU-accelerated packet classification
2. bloom search algorithm
3. GSwitch

GPGPU - General Purpose GPU programming

- OpenCL, CUDA
- each thread executes the same program (*kernel*)
- utilization and performance relies on multiple factors
 - synchronization
 - memory access patterns (coalesced reads)
 - caching (on-chip SRAM, shared memory)
 - pipelining (copy & execution)
 - work allocation

GPGPU - GPU architecture



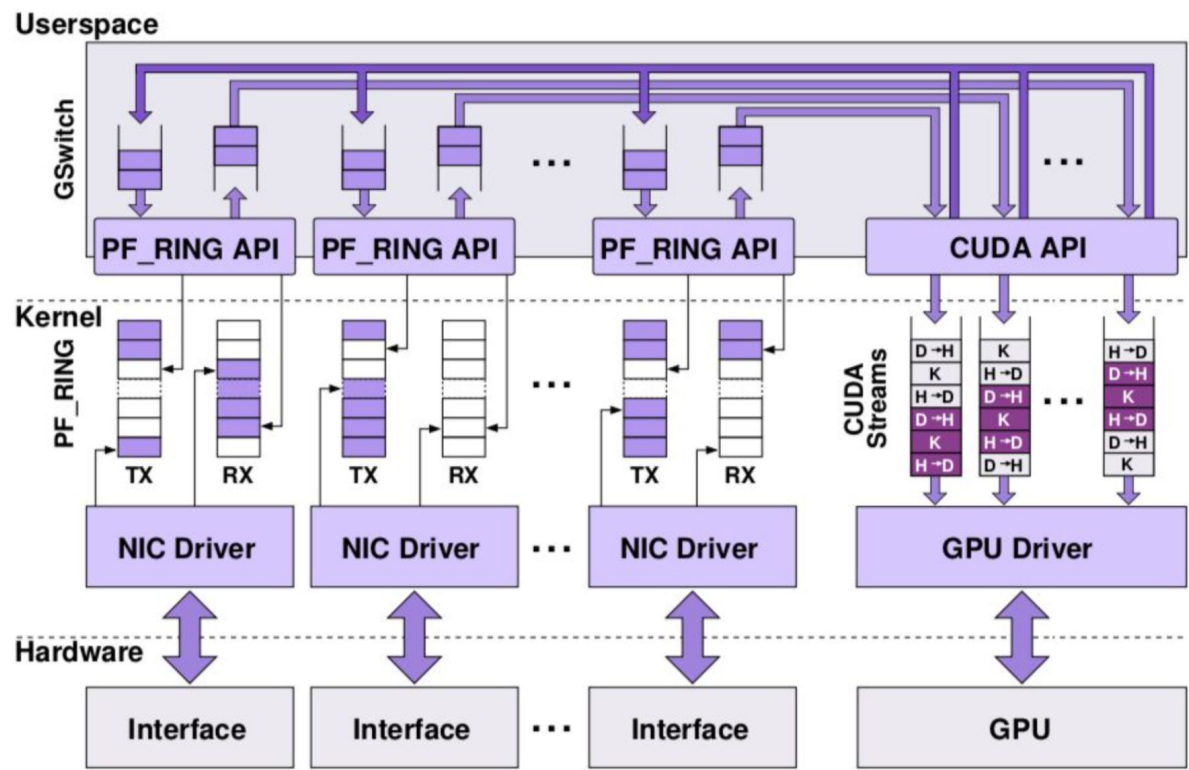
GPGPU - implemented algorithms

- linear search
 - go through all rules
- tuple search
 - divide rules into classes defined by a mask
 - lookup in each class hash table
- bloom search

bloom filters

- probabilistic data structure
- quickly (constant time) tell if element is not in set
- might get false positives, rate is tunable
- bloom search
 - quickly decide if need to check class based on filter

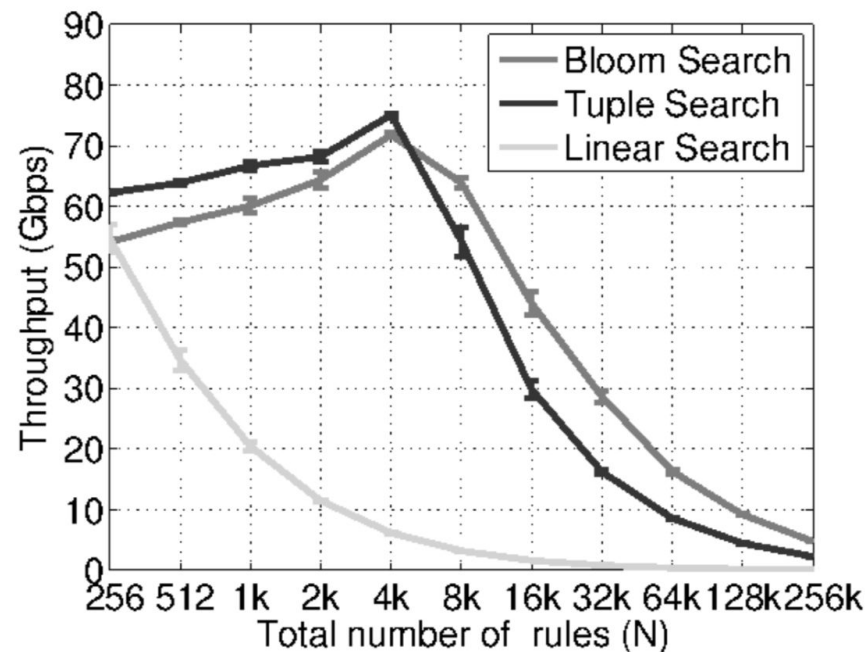
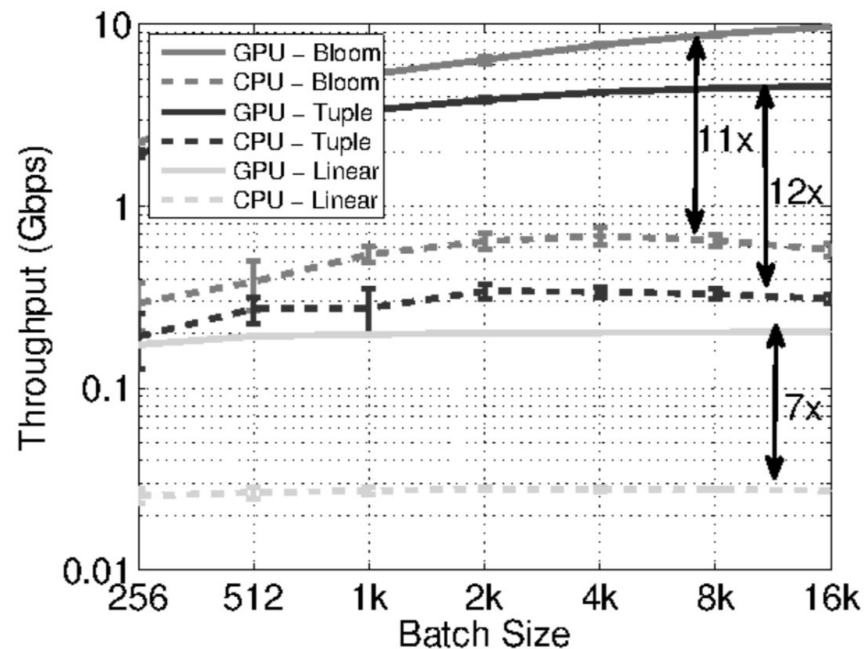
GSwitch



results and evaluation

- single CPU, single GPU of comparable price
- GPU micro-benchmarks to tune parameters
- speedup: 7x (linear search), 11x (tuple search), 12x (bloom search)
- “64-byte frames at 30 Gbps and a maximum per-packet latency of 500 μ s”.
- shift performance bottleneck from packet classification back to packet I/O

results and evaluation




Outline

- GPU-based packet classification [1]
 - packet classification overview
 - problem statement
 - GPU programming
 - bloom filters
 - GSwitch
 - results & evaluation
- **safe network updates [3]**
 - problem statement
 - inconsistent update example
 - update mechanisms
 - mathematical model
 - results & evaluation

problem statement

- computer networks are complex
- updates involve changing configurations on several nodes
- updates often introduce incorrect transient states
- instability, outages, performance disruptions, security vulnerabilities

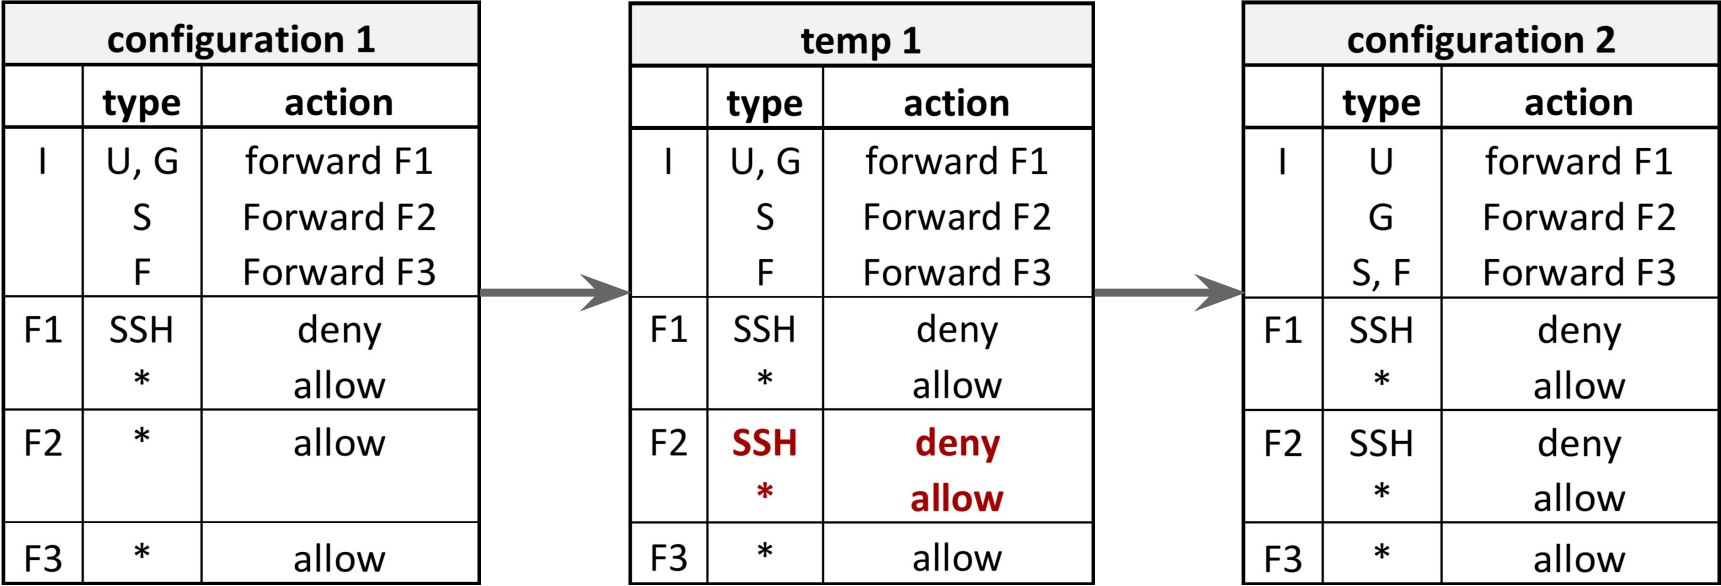
inconsistent updates - example



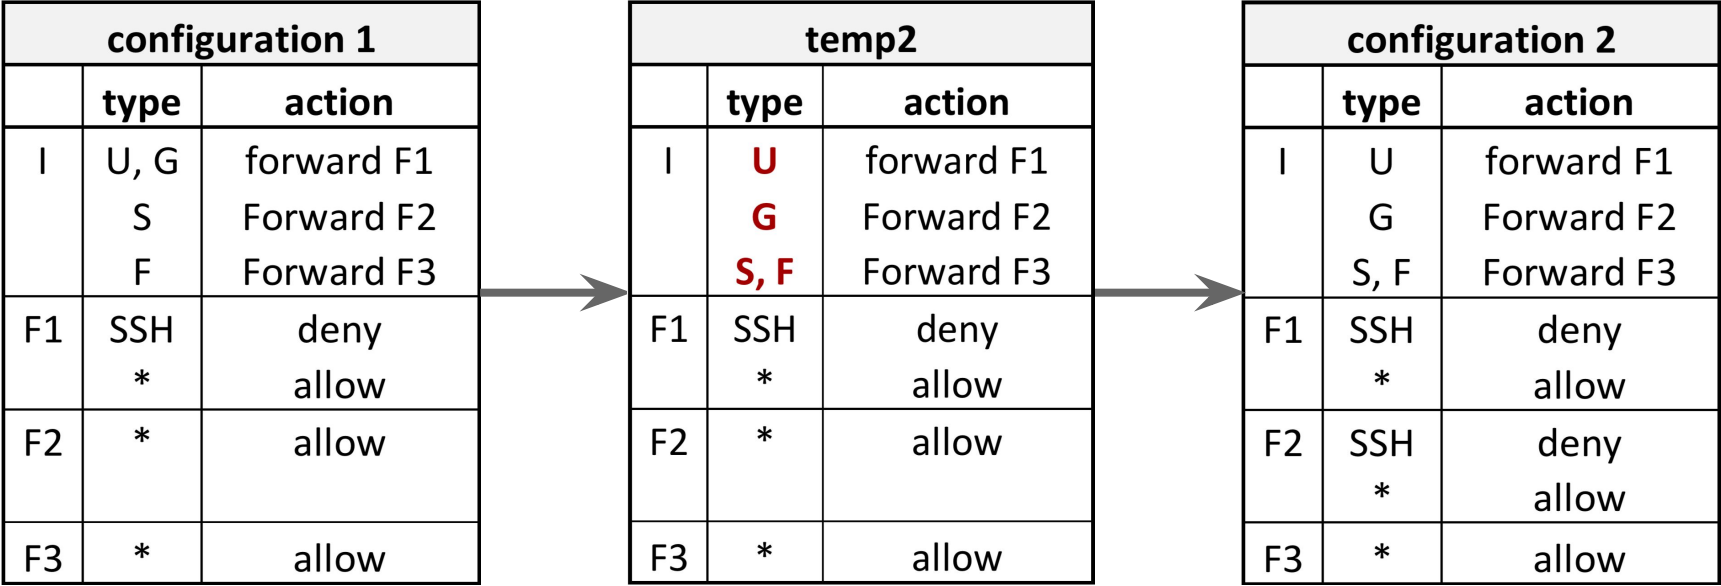
configuration 1		
	type	action
I	U, G S F	forward F1 Forward F2 Forward F3
F1	SSH *	deny allow
F2	*	allow
F3	*	allow

configuration 2		
	type	action
I	U G S, F	forward F1 Forward F2 Forward F3
F1	SSH *	deny allow
F2	SSH *	deny allow
F3	*	allow

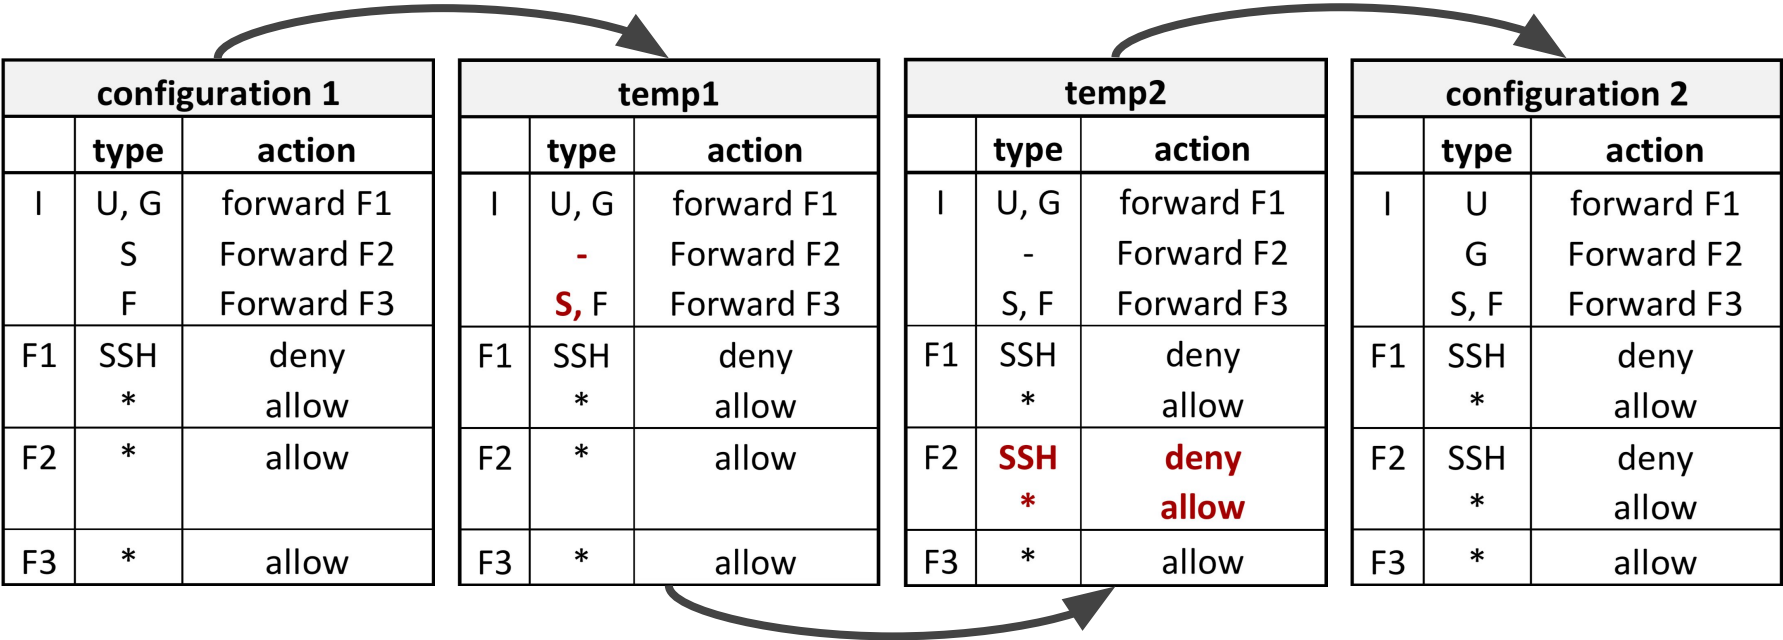
inconsistent update 1.



inconsistent update 2.



correct update



goal: consistent, safe network updates

- network abstraction
 - developers should not need to design step-by-step network update
 - offer high-level primitives and let the system make sure it is correct
- consistency levels
 - per-packet consistency (vs. atomic)
 - per-flow consistency (e.g. HTTP load balancing)

update mechanisms

- **2-phase update** (per-packet consistent)
 1. pre-install new config on all nodes with version number (*unobservable update*)
 2. unlock new policy on ingress ports by stamping packets (*one-touch update*)
- **switch rules with timeouts** (per-flow consistency)
 1. pre-install new config on all nodes
 2. install new config on ingress switches with low priority
 3. set soft timeout on the old config

mathematical modeling

Bit	$b ::= 0 \mid 1$
Packet	$pk ::= [b_1, \dots, b_k]$
Port	$p ::= 1 \mid \dots \mid k \mid Drop \mid World$
Located Pkt	$lp ::= (p, pk)$
Trace	$t ::= [lp_1, \dots, lp_n]$
Update	$u \in LocatedPkt \rightarrow LocatedPkt \text{ list}$
Switch Func.	$S \in LocatedPkt \rightarrow LocatedPkt \text{ list}$
Topology Func.	$T \in Port \rightarrow Port$
Port Queue	$Q \in Port \rightarrow (Packet \times Trace) \text{ list}$
Configuration	$C ::= (S, T)$
Network State	$N ::= (Q, C)$

results and evaluation

- implemented on top of OpenFlow/Kinetic
 - simple interface: `per_packet_update`, `per_flow_update`
 - automatically implement transition based on topology, config, consistency level
 - store version tag in VLAN field
- Mininet network simulation (fat-tree, small-world, random topologies)
 - routing, multicast test scenarios
 - 20-100% overhead (in terms of rules installed)

References

- [1] Varvello, Matteo et al. “Multi-Layer Packet Classification with Graphics Processing Units.” CoNEXT (2014).
- [2] Taylor, D. E.. “Survey and taxonomy of packet classification techniques.” ACM Comput. Surv. 37 (2005).
- [3] Reitblatt, Mark et al. “Abstractions for network update.” SIGCOMM (2012).

Thank you!